



UX & Journey-Management Review

Usability & experience recommendations for the v5 platform mockup, ahead of production release.

PRODUCT	Intendra AI — customer-behaviour intelligence platform (omnichannel intent, root-cause, best-fit, coaching)
ARTEFACT	intendra-app-v5.html · 11 views reviewed on desktop & at mobile width
LENS	Senior UX + CX journey-management; Nielsen heuristics, WCAG 2.1 AA, value-irritant (Bill Price)
DATE	13 August 2026
STATUS	Advisory — no changes made to the mockup

Verdict — a strong, differentiated product with an unusually coherent narrative structure. It is not yet production-ready on three fronts: responsive support, accessibility/contrast, and first-time comprehension. None are structural; all are addressable before ship.

EXECUTIVE SUMMARY

A rare thing: a CX analytics product with a point of view

Intendra reads as a single argument rather than a pile of dashboards. The primary navigation is a narrative — **what happened** → **why customers reached out** → **how we served** → **who served** → **who fits what** → **the evidence** → **what to do** — and every number traces back to a scored conversation. It applies the value–irritant framework as concrete, costed actions, governs its definitions in one place (Studio), and closes the loop from initiative to expected-vs-realised dollars. That spine is the product's biggest asset and should be protected as it hardens for production.

The gaps are not conceptual — they are the finishing layer that enterprise buyers scrutinise. The interface is desktop-only and breaks at tablet/phone widths; secondary text and color-only signals fall below accessibility thresholds that procurement (VPAT/Section 508) will test; and the dense, jargon-rich surface assumes an expert reader with no onboarding scaffold. The recommendations below are grouped into eight themes and tagged by priority, followed by a visual-design-system section covering colour, type and shape. Fixing the P0 items is the shortest path to "solid and easy to use."

THEME	READINESS	HEADLINE
Information architecture	Strong	Narrative nav is a genuine differentiator — preserve it
Responsive / device support	At risk	Layout does not reflow; nav truncates on tablet & phone
Accessibility & contrast	At risk	Grey-on-grey text and color-only signals likely fail AA
Visual hierarchy & rhythm	Good	Empty vertical bands; non-sticky table headers
Data visualisation	Good	Powerful but dense; truncated labels, inline formulae
Journey management depth	Opportunity	Contact-centric; elevate the cross-contact journey
Content & onboarding	Needs work	Heavy jargon (NEM, DIHT, snowballs) with no glossary
Interaction & affordance	Good	Clickable cards look static; confirm filter persistence
Visual design system	Good	Real identity, but untokenised — 29 type sizes, 13 radii

Priority key — **P0** blocks a confident production release · **P1** materially hurts usability, fix in the first hardening sprint · **P2** polish that raises perceived quality.

PRESERVE THESE

What is already working — and must survive the redesign

It is easy for a hardening pass to sand off the very things that make this product distinctive. Treat the following as fixed points.

- **Narrative navigation.** The tab order tells a story and gives users a durable mental model of where answers live. Most competitors present unordered dashboards; this is a real advantage.
- **One denominator, one source.** 8,470 contacts holds steady across every view, and Studio governs the definitions with a change log and second-admin approval. This is the foundation of analyst trust.
- **Value–irritant made actionable.** Every intent carries a costed disposition (Digitise / Leverage / Eliminate) and an "if fixed" dollar figure — Bill Price's framework turned into a work queue.
- **Closed loop.** Actions ties each initiative to an owner, an expected path, and realised-vs-expected dollars — the discipline most CX tools never reach.
- **Evidence at the bottom of every claim.** Conversations shows the transcript with inline signal tags and an intent → root-cause → outcome chain. Nothing is asserted without a way to inspect it.
- **Editorial voice.** Plain-language sentences ("The largest single source of defects...") make dense analytics legible. Keep the human tone.

Scope & method

All eleven views were walked end-to-end on desktop (Summit, Demand, Quality, People, Fit, Conversations, Actions, Metrics, Reports, Studio, Ask Intendra) and re-checked at a 390 px mobile width. This is a heuristic expert review of a static mockup: it evaluates structure, layout, legibility, and journey logic. It does not replace moderated usability testing with real CX leads and QA managers, nor a formal accessibility audit — both are recommended before or shortly after launch and are noted where relevant.

RECOMMENDATIONS

Findings by theme

Each card states the finding, why it matters, and the recommended fix. Themes A–H; priority tag at right.

A · Responsive & device support

A1 The layout does not reflow below desktop

P0 · CRITICAL

- FINDING

At tablet and phone widths the KPI cards keep their fixed width and overflow sideways (the whole page gets a horizontal scrollbar), and the top navigation truncates — Reports, Studio and Ask Intendra scroll off-screen with no menu to reach them.
- IMPACT

CX directors and QA managers increasingly review these numbers on a tablet or phone between meetings. Today those users cannot reach four of the eleven views and must pan horizontally to read a single card row.
- FIX

Define breakpoints before ship: stack KPI cards to 2-up / 1-up, collapse the nav into a menu (or priority-plus pattern) below ~1024 px, and constrain wide tables to scroll inside their own container rather than the page body.

A2 Wide data tables scroll the whole page horizontally

P1 · HIGH

- FINDING

The Demand and People tables are wide; on smaller viewports the entire page scrolls sideways instead of the table alone.
- IMPACT

Horizontal page scroll is disorienting and hides the navigation and filters as the user reads across columns.
- FIX

Wrap each table in an overflow-x container with the first (label) column pinned, so the row header stays visible while the metrics scroll.

B · Accessibility & contrast

B1 Secondary text is likely below WCAG AA contrast**P0 · CRITICAL**

FINDING Captions, micro-labels (the 8 px uppercase greys), "12 wk" annotations and much table meta are light grey on the light-grey (#f4f6f4) page. Several of these pairings appear to fall under the 4.5:1 ratio AA requires for body text.

IMPACT Beyond readability, enterprise procurement routinely requires a VPAT / Section 508 statement; contrast failures are the most common blocker and can stall a deal after the demo has won.

FIX Run every text/background pair through a contrast checker and darken secondary text to $\geq 4.5:1$ ($\geq 3:1$ for large text). Budget a formal WCAG 2.1 AA audit into the release plan.

B2 Meaning is carried by color alone**P1 · HIGH**

FINDING Red vs grey dots in the People and Fit bee-swarms, red numbers, and green vs grey bars encode status purely by hue. Roughly 8% of male users have a color-vision deficiency.

IMPACT Color-blind users cannot distinguish "below target" from "on target," and the signal is lost in grayscale printing of the very reports this product generates.

FIX Add a redundant non-color cue everywhere status is shown — a shape, an icon, a border, or a short text label (e.g. "below 85") alongside the color.

B3 Green is overloaded — brand, links, and "good," but also "cost"**P1 · HIGH**

FINDING Green is the brand accent and link color and the "positive" signal — yet on the Demand Pareto the green bar segment represents *downstream cost*, a negative. The same hue means opposite things.

IMPACT Users read green as good; encoding a cost you want to reduce in green quietly inverts the intuitive signal and erodes at-a-glance trust.

FIX Reserve green for brand/interaction and genuine positives. Give cost and downstream their own neutral or warm hue, and document a small semantic palette (positive / negative / neutral / brand) applied consistently.

C · Visual hierarchy & layout rhythm

C1 Large empty vertical bands strand the eye**P1 · HIGH**

FINDING On Summit (between "Where it concentrates" and "What's happening"), and around the tall charts on Fit and Actions, unbalanced two-column sections leave big blank regions the reader must scroll past.

IMPACT The whitespace reads as "nothing here," breaks the sense of a continuous story, and makes pages feel longer than their content.

FIX Balance the paired columns (match heights or let the shorter column carry a summary), and tighten the vertical spacing scale so sections follow each other without dead air.

C2 Table headers are not sticky**P1 · HIGH**

FINDING The Demand intent table and the Quality parameter table run ~20 rows; their column headers (Contacts, Cost, Downstream, Action...) scroll away, leaving rows of unlabelled numbers.

IMPACT Users lose the meaning of columns mid-scan and must scroll back up to re-orient — a repeated small tax on the product's densest, most-used screens.

FIX Pin the header row (and ideally the section's summary line) while the table body scrolls.

C3 Content width is inconsistent between pages**P2 · POLISH**

FINDING Most views use a wide container (~1400 px) while Reports sits noticeably narrower (~1160 px).

IMPACT Shifting margins between tabs read as a small wobble in an otherwise very controlled design system.

FIX Standardise one max content width (with documented exceptions) across all routes.

D · Data visualisation

D1 Rotated, truncated category labels on the Pareto**P1 · HIGH**

FINDING The Demand Pareto rotates x-axis labels and truncates them ("Get a delivery...", "Update my payment..."), and overlays stacked bars, a cumulative line, and a 57%/80% reference at once.

IMPACT Truncated diagonal labels are slow to read and the triple-encoding is a lot to decode in the product's entry point to "why customers contact us."

FIX Prefer horizontal bars for long intent names (labels read straight and in full), and consider letting the cumulative line / reference be a toggle so the first read is simpler.

D2 Glanceable charts have no scale or on-demand value**P2 · POLISH**

FINDING Sparklines and the 12-week bar strips carry no axis or values; a trend with no scale can imply a bigger move than the data supports.

IMPACT Fine as a glance, but users will want the number when a line looks alarming — and can't get it.

FIX Add a hover/tap tooltip that reveals the value and period on every spark and mini-bar.

D3 Scoring formulae are shown inline, competing with the chart**P2 · POLISH**

FINDING Fit's excellent performance quadrant sits beside the full "how the axes are built" weighting recipe, and the Demand rows show unexplained $\times 0.5$ / $\times 1$ / $\times 0.7$ multiplier badges.

IMPACT First read should be the picture; the methodology is reference material most users want only occasionally.

FIX Move the weighting recipe and multiplier definitions into a collapsed "How this is scored" disclosure and a hover legend, keeping the chart as the hero.

E · Journey-management depth

E1 The product is contact-centric; the journey is a footnote**P1 · HIGH**

FINDING	Cross-contact reality ("3 contacts in 90 days") appears only inside a single conversation. Repeat/rework chains, channel-switching, and effort accumulated across a customer's journey are not a first-class lens.
IMPACT	Journey managers — a core buyer for this product — think in end-to-end journeys, not single contacts. Repeat contact and channel-switching are where avoidable cost and churn actually compound.
FIX	Elevate the thread to a first-class Journey view: repeat-contact chains, channel hops, and cumulative effort per resolved need, tied back to the intents and root causes already detected.

E2 Prioritisation uses value but not effort**P2 · POLISH**

FINDING	"What to act on" ranks by dollars at stake — strong. But it omits the effort/feasibility axis the value–irritant quadrant elsewhere in the product already implies.
IMPACT	Teams need quick wins, not just the biggest prize; a \$340k fix that needs a carrier contract may rank below a \$90k fix shippable this sprint.
FIX	Add an effort/feasibility estimate and offer a value-vs-effort view so the shelf surfaces quick wins as well as big bets.

E3 Best-fit is analytical but not yet actionable**P2 · POLISH**

FINDING	Fit identifies the best agent↔intent↔channel pairing and hints at "route away, coach" — but the analysis does not connect to a routing or coaching action.
IMPACT	The insight's value is only realised when it changes routing or a coaching plan; today that hand-off lives outside the tool.
FIX	From Fit, offer a one-click hand-off — "route Returns to this pool," "add these 6 to a coaching plan" — even if it only drafts the task in the connected system.

F · Content & onboarding

F1 Expert jargon with no glossary or definitions**P0 · CRITICAL**

- FINDING** NEM, CES, DIHT, "snowballs melted," the Stars/Anchors/Sleepers/Weak quadrant, "auto-zero the audit," and "capability-cluster" all appear without inline explanation. "Snowballs" in particular is opaque on first read.
- IMPACT** A first-time CX lead or QA manager cannot self-serve; the product assumes training. This is the single biggest barrier to "easy to use" for anyone outside the design team.
- FIX** Add hover/tap definitions on every coined term and a lightweight glossary, plus a one-line "what this page answers" descriptor under each title (the aria titles already written are a perfect source).

F2 Navigation labels don't explain themselves**P2 · POLISH**

- FINDING** Terse tabs ("Fit," "People") are elegant but unlabelled; the beautiful full-sentence descriptions exist only as hidden hover/aria text.
- IMPACT** New users can't predict what "Fit" contains, so they click around to build a map that the copy could give them for free.
- FIX** Surface the descriptor as a visible tooltip on hover, or as a persistent sub-label, so wayfinding is legible from the first session.

F3 Empty and first-run states are undefined**P2 · POLISH**

- FINDING** Ask Intendra models this well with suggested prompts; other views assume a full, mature dataset.
- IMPACT** A brand-new tenant, a filtered-to-empty view, or a partner with no data will hit blank regions with no guidance.
- FIX** Design explicit empty states ("no fatal violations this week," "not enough data to score this pool yet") that reassure rather than look broken.

G · Interaction, affordance & state

G1 Drill-downs are discoverable only by trial**P1 · HIGH**

- FINDING** KPI cards, table rows and chart bars are clickable ("Open details") but look static — no hover elevation or cursor affordance signals the depth beneath them.
- IMPACT** The product's best feature — that every number opens into its evidence — is easy to miss, so users under-use the drill-downs that set it apart.
- FIX** Add clear interactive affordances: hover elevation/underline, a pointer cursor, and a subtle "open" chevron on anything that drills in.

G2 Confirm that global filters persist across views**P1 · HIGH**

- FINDING** Channel / period / partner / team filters repeat on each page. It must be verified that filtering to, say, "Voice · Partner MK" holds as the user moves Summit → Quality → People.
- IMPACT** If filters reset per view, an analyst re-sets context on every tab and can unknowingly compare a filtered view against an unfiltered one — a correctness risk, not just an annoyance.
- FIX** Make filter state global and visibly persistent, with a single obvious "clear all" and a breadcrumb of the active scope.

G3 State-changing actions need confirmation and undo**P2 · POLISH**

- FINDING** "Create initiative," "Adopt from the shelf" and alert-rule edits change program state directly from a list.
- IMPACT** Accidental adoption or an errant rule change is easy from a dense list and there's no visible safety net.
- FIX** Add a lightweight confirm (or an undo toast) on actions that write state, consistent with the change-log discipline Studio already models.

VISUAL DESIGN SYSTEM

Colour, type & shape — rationalising what's already there

The mockup is not visually undesigned — it has a genuine identity: an editorial **Georgia** serif for display headlines, a system **sans** for the interface, a **monospace** for the uppercase labels, a vivid **mint** brand green, a deep-green interaction/positive green, and a warm-**gold** highlight. The goal here is not to restyle it but to **tokenise** it — collapse the accidental variation into a small, named set of colours, sizes and shapes so the identity holds together as the product grows. All values below are read directly from the v5 file.

1 · Colour

The palette as built (CSS custom properties in the file):

#1bea87 --brand	#0f7d51 --accent / --pos	#e7f8ef --accent-soft
#c03d33 --red / --neg	#faedeb --red-soft	#c9a23f --hlb
#f6efdc --hl	#28282d --ink	#3f434b --body
#6d717b --muted	#767b85 --faint	#e4e5e9 --line
#f5f6f8 --bg	#ffffff --card	

Contrast audit — measured against the page/card background (WCAG 2.1 AA needs 4.5:1 for normal text, 3:1 for large):

TOKEN	RATIO	VERDICT (small text)	USE IT FOR
--ink #28282d	13.6:1	Pass	Headings — no change
--body #3f434b	9.2:1	Pass	Body — no change
--muted #6d717b	4.5:1	Pass (just)	Secondary — hold the line, don't lighten
--faint #767b85	3.9:1	FAIL	Captions/labels — darken to ~#5b606b
--accent #0f7d51	4.8:1	Pass	Links, positive text
--red #c03d33	4.9:1	Pass	Alerts, negative text
--brand #1bea87	1.6:1	FAIL	Fills & logo ONLY — never text/thin lines
--hlb gold #c9a23f	2.4:1	FAIL	Borders/fills ONLY — never text

Colour recommendations

- **Darken** - - faint. At 3.9:1 it fails AA for the small captions and 8 px labels it's used on. Move it to roughly **#5b606b** (≈5.5:1) and reserve today's lighter grey for large text only. **[P0 · ties to finding B1]**
- **Never set text or thin marks in mint or gold.** - - brand #1bea87 (1.6:1) and - - hlb #c9a23f (2.4:1) are for fills, logos and borders only — a mint sparkline or gold caption is effectively invisible.
- **Separate the two greens by role and stick to it.** Mint = identity/brand energy; deep green = interaction + positive. They currently blur — e.g. the mint fill also encodes *downstream cost* on Demand, which reads as "good." **[P1 · ties to B3]**
- **Give cost/negative-magnitude its own swatch.** Introduce a neutral-warm hue (e.g. slate or clay) for "cost" and "downstream" so green is freed to mean only positive.
- **Formalise four semantic pairs.** positive (#0f7d51 / #e7f8ef), negative (#c03d33 / #faedeb), attention (gold #c9a23f / #f6efdc), neutral (grey ramp) — each a text/fill pair — and keep - - brand out of the semantic set entirely.

2 · Typography

Keep the three-family system — Georgia serif for display, system-sans for UI, mono for labels. The serif headline is a real editorial signature and worth protecting; few competitors have a voice this distinct. The problem is not the fonts, it's the **lack of a scale**.

- The file uses **~29 distinct font sizes**, many at half-pixel values (12.5, 11.5, 9.5, 13.5, 14.5, 15.5 px). There is no modular rhythm, so headings and labels drift a pixel apart between components.
- The **body base is 12.5 px** — small for a data-dense enterprise tool viewed for long stretches; it compounds the contrast issue on secondary text.
- Uppercase label **tracking is set at ~10 different values** (.03–.14em), so the kicker labels don't feel uniform.

Proposed 8-step scale (round to whole pixels, ~1.25 ratio):

STEP	SIZE	FAMILY / WEIGHT	USE
Display	28 px	Georgia serif	Page title ("How the week went")
H1	20 px	Sans bold	Section headers
H2	16 px	Sans bold	Sub-sections, card titles
Body-L	14 px	Sans regular	Lead sentences, emphasis
Body	13 px	Sans regular	Default body — up from 12.5
Small	12 px	Sans regular	Secondary / table meta
Caption	11 px	Sans regular	Captions (with darkened grey)
Micro	10 px	Mono, +0.1em	Uppercase kicker labels only

- **Snap every size to the 8 steps** and delete the half-pixel values — a single source of truth kills the drift. **[P1]**
- **Lift body to 13 px** and reserve 10–11 px for genuine labels, not readable content.
- **Collapse label tracking to one value** (0.1em) and use the mono face consistently for those uppercase kickers.
- **Keep the serif for page titles only** — using it lower down would dilute the editorial signal.

3 · Shape, radius & spacing

The mockup uses **13 different corner radii** — 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 14 px plus circles and asymmetric bubbles. Cards appear at 12, 10, 8 *and* 6 px in different places; pills and chips vary between 3 and 7 px. The eye reads this as slight sloppiness in an otherwise controlled design. Collapse to a **three-tier radius token set**:

TOKEN	VALUE	APPLIES TO
radius-sm	4 px	Tags, chips, inputs, small toggles
radius-md	8 px	Buttons, segmented controls, table & chart wrappers
radius-lg	12 px	Cards, panels, modals, KPI tiles
radius-pill	999 px	Status pills, filter toggles
radius-circle	50%	Dots, avatars, agent markers

- **Map every element to one of the three radii.** Keep the one deliberate exception — the asymmetric chat bubble (12/12/3/12) — since it signals "message." [P2]
- **Adopt a 4-px spacing grid** (4 / 8 / 12 / 16 / 24 / 32 / 48). The stranded empty bands in finding C1 come partly from ad-hoc gaps; a grid removes the guesswork and the dead air.
- **Set minimum sizes for the responsive build:** ≥44 px touch targets, ≥1 px hairlines that survive scaling, and consistent 1-px - - line borders rather than mixed weights.
- **Standardise elevation.** Cards currently rely on borders alone; define one or two shadow tokens (e.g. a resting card and a hover-lift) so the drill-down affordance in finding G1 has a consistent physical language.
- **Unify iconography** to a single stroke weight and size grid (16 / 20 px) — the arrow, chevron and status glyphs should feel like one set.

Why tokenise at all

None of this changes how Intendra looks in a single screenshot — it changes whether the 50th screen a developer builds still matches the first. A named set of colour, type and shape tokens is what keeps the strong identity you already have from fraying as the product scales, and it is the layer an accessibility audit and a responsive rebuild both depend on. Do it now, while the surface is still small.

BEFORE YOU SHIP

The P0 shortlist

Seven items stand between the current mockup and a confident production release. Everything else raises quality; these remove blockers.

P0 — release blockers

- | | |
|------|---|
| ■ A1 | Make the layout responsive — reflow cards, collapse the nav, contain wide tables below desktop. |
| ■ B1 | Fix secondary-text contrast to WCAG AA and schedule a formal accessibility audit (VPAT/508). |
| ■ F1 | Add inline definitions + a glossary so a first-time user can self-serve the jargon. |

P1 — first hardening sprint

- | | |
|-----------|---|
| ■ B2 / B3 | Remove color-only signals; stop encoding cost in green — adopt a documented semantic palette. |
| ■ C1 / C2 | Close the empty vertical bands and pin table headers on the dense screens. |
| ■ D1 | Switch the Demand Pareto to horizontal bars / full labels; simplify the default read. |
| ■ E1 | Elevate the cross-contact journey to a first-class lens. |
| ■ G1 / G2 | Signal that cards drill in; verify global filters persist across views. |
| ■ Design | Tokenise the visual system: 8-step type scale, 3 radii, 4 semantic colour pairs. |

A closing note

The hard, original work in Intendra is already done and it is genuinely good — the narrative spine, the value–irritant discipline, the traceable evidence, the closed loop. The remaining work is the finishing layer: make it reach every device, make it legible to every user and every eye, and let newcomers learn it without a trainer. Do those, and "solid, easy-to-use product" is well within reach for this release.

Recommended next step: one round of moderated usability testing with 5–6 real CX leads and QA managers against the P0/P1 fixes, to confirm the jargon and journey changes land before the wider rollout.